

# Experiment 11: Heart Disease Prediction using different Machine Learning Models

```
In [ ]: #libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score, classification_report
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import (
    confusion_matrix, classification_report, roc_curve, roc_auc_score
)
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [74]: # Load the dataset
df = pd.read_csv('heart_disease.csv')
df
```

```
Out[74]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	sl
<b>0</b>	52	1	0	125	212	0	1	168	0	1.0	
<b>1</b>	53	1	0	140	203	1	0	155	1	3.1	
<b>2</b>	70	1	0	145	174	0	1	125	1	2.6	
<b>3</b>	61	1	0	148	203	0	1	161	0	0.0	
<b>4</b>	62	0	0	138	294	1	1	106	0	1.9	
...	...	...	...	...	...	...	...	...	...	...	...
<b>1020</b>	59	1	1	140	221	0	1	164	1	0.0	
<b>1021</b>	60	1	0	125	258	0	0	141	1	2.8	
<b>1022</b>	47	1	0	110	275	0	0	118	1	1.0	
<b>1023</b>	50	0	0	110	254	0	0	159	0	0.0	
<b>1024</b>	54	1	0	120	188	0	1	113	0	1.4	

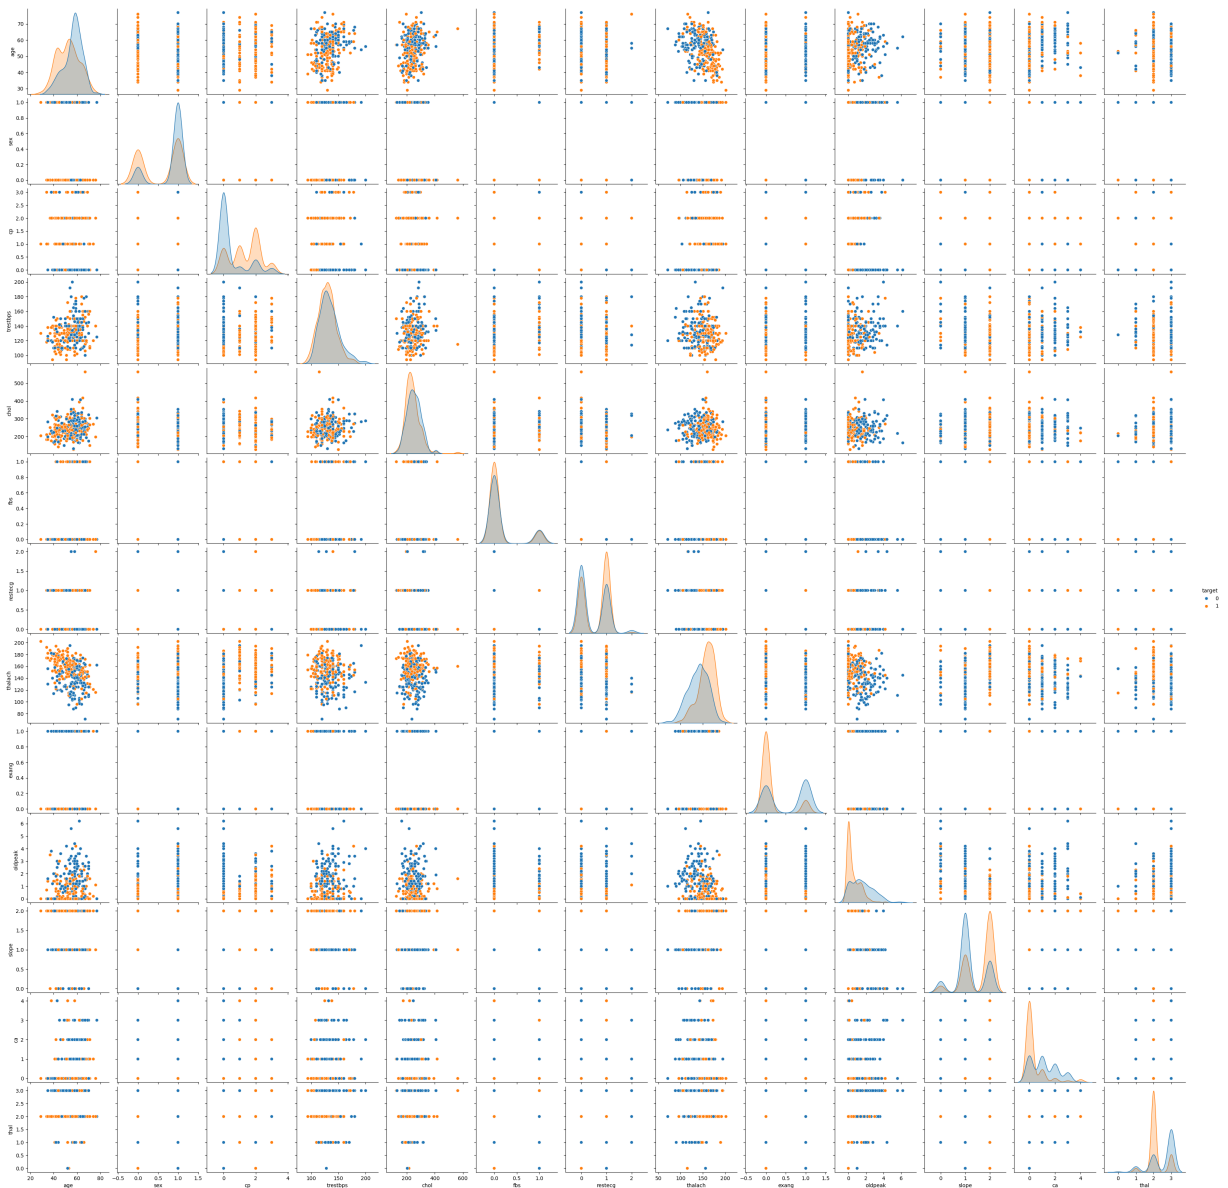
1025 rows × 14 columns

```
In [75]: df.describe()
```

```
Out[75]:
```

	age	sex	cp	trestbps	chol	sl
<b>count</b>	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
<b>mean</b>	54.434146	0.695610	0.942439	131.611707	246.000000	0.100000
<b>std</b>	9.072290	0.460373	1.029641	17.516718	51.59251	0.300000
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000
<b>25%</b>	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000
<b>50%</b>	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

```
In [76]: sns.pairplot(df, hue='target')
plt.show()
```



```
In [77]: # Split the data into features (X) and target variable (y)
X = df.drop('target', axis=1)
y = df['target']
```

```
In [78]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [79]: # Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [105]: models = {
    'Logistic Regression': LogisticRegression(solver='liblinear'),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'SVM': SVC(probability=True),
    'KNN': KNeighborsClassifier(),
    'Naive Bayes': GaussianNB(),
```

```
}
```

```
for name, model in models.items():  
    print(f"Training {name}...")  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
    accuracy = accuracy_score(y_test, y_pred)  
    print(f"{name} Accuracy: {accuracy:.2f}")  
    print(classification_report(y_test, y_pred))
```

Training Logistic Regression...

Logistic Regression Accuracy: 0.81

	precision	recall	f1-score	support
0	0.86	0.75	0.80	159
1	0.76	0.87	0.81	149
accuracy			0.81	308
macro avg	0.81	0.81	0.80	308
weighted avg	0.81	0.81	0.80	308

Training Decision Tree...

Decision Tree Accuracy: 0.97

	precision	recall	f1-score	support
0	0.95	1.00	0.97	159
1	1.00	0.94	0.97	149
accuracy			0.97	308
macro avg	0.97	0.97	0.97	308
weighted avg	0.97	0.97	0.97	308

Training Random Forest...

Random Forest Accuracy: 0.98

	precision	recall	f1-score	support
0	0.96	1.00	0.98	159
1	1.00	0.96	0.98	149
accuracy			0.98	308
macro avg	0.98	0.98	0.98	308
weighted avg	0.98	0.98	0.98	308

Training SVM...

SVM Accuracy: 0.89

	precision	recall	f1-score	support
0	0.92	0.86	0.89	159
1	0.86	0.92	0.89	149
accuracy			0.89	308
macro avg	0.89	0.89	0.89	308
weighted avg	0.89	0.89	0.89	308

Training KNN...

KNN Accuracy: 0.84

	precision	recall	f1-score	support
0	0.87	0.82	0.84	159
1	0.82	0.87	0.84	149
accuracy			0.84	308
macro avg	0.84	0.84	0.84	308
weighted avg	0.84	0.84	0.84	308

Training Naive Bayes...

Naive Bayes Accuracy: 0.81

	precision	recall	f1-score	support
0	0.88	0.74	0.81	159
1	0.76	0.89	0.82	149
accuracy			0.81	308
macro avg	0.82	0.82	0.81	308
weighted avg	0.82	0.81	0.81	308

```
In [85]: for name, model in models.items():
          print(f"Training {name}...")
          model.fit(X_train, y_train)
          y_pred = model.predict(X_test)
          accuracy = accuracy_score(y_test, y_pred)
          print(f"{name} Accuracy: {accuracy:.2f}")
          print(classification_report(y_test, y_pred))
          print("\n")
```

Training Logistic Regression...

Logistic Regression Accuracy: 0.81

	precision	recall	f1-score	support
0	0.86	0.75	0.80	159
1	0.76	0.87	0.81	149
accuracy			0.81	308
macro avg	0.81	0.81	0.80	308
weighted avg	0.81	0.81	0.80	308

Training Decision Tree...

Decision Tree Accuracy: 0.97

	precision	recall	f1-score	support
0	0.95	1.00	0.97	159
1	1.00	0.94	0.97	149
accuracy			0.97	308
macro avg	0.97	0.97	0.97	308
weighted avg	0.97	0.97	0.97	308

Training Random Forest...

Random Forest Accuracy: 1.00

	precision	recall	f1-score	support
0	1.00	1.00	1.00	159
1	1.00	1.00	1.00	149
accuracy			1.00	308
macro avg	1.00	1.00	1.00	308
weighted avg	1.00	1.00	1.00	308

Training SVM...

SVM Accuracy: 0.89

	precision	recall	f1-score	support
0	0.92	0.86	0.89	159
1	0.86	0.92	0.89	149
accuracy			0.89	308
macro avg	0.89	0.89	0.89	308
weighted avg	0.89	0.89	0.89	308

Training KNN...

KNN Accuracy: 0.84

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.87	0.82	0.84	159
	1	0.82	0.87	0.84	149
accuracy				0.84	308
macro avg		0.84	0.84	0.84	308
weighted avg		0.84	0.84	0.84	308

Training Naive Bayes...

Naive Bayes Accuracy: 0.81

		precision	recall	f1-score	support
	0	0.88	0.74	0.81	159
	1	0.76	0.89	0.82	149
accuracy				0.81	308
macro avg		0.82	0.82	0.81	308
weighted avg		0.82	0.81	0.81	308

```
In [121]: # Definitions for the accuracy variables
LRA = 0.81 # Logistic Regression accuracy
KNNA = 0.84 # K-Nearest Neighbour accuracy
SVMA = 0.89 # Support Vector Machine accuracy
GNBA = 0.81 # Gaussian Naive Bayes accuracy
DTA = 0.97 # Decision Tree accuracy
RFA = 1.00 # Random Forest accuracy

# Create the DataFrame with defined accuracy values
compare = pd.DataFrame({
    'Model': ['Logistic Regression', 'K-Nearest Neighbour', 'Support Vector Machine', 'Gaussian Naive Bayes', 'Decision Tree', 'Random Forest'],
    'Accuracy': [LRA * 100, KNNA * 100, SVMA * 100, GNBA * 100, DTA * 100, RFA * 100]
})

# --- Create Accuracy Comparison Table ---
styled_compare = compare.style.set_properties(**{'font-family': 'Segoe UI'})
styled_compare
```

Out[121]:

Model	Accuracy
Logistic Regression	81.000000
K-Nearest Neighbour	84.000000
Support Vector Machine	89.000000
Gaussian Naive Bayes	81.000000
Decision Tree	97.000000
Random Forest	100.000000



```

In [111]: # Dictionary for model abbreviations
model_abbreviations = {
    'Decision Tree': 'DT',
    'Logistic Regression': 'LR',
    'Random Forest': 'RF',
    'SVM': 'SVM',
    'KNN': 'KNN',
    'Naive Bayes': 'NB'
}

# Plot ROC Curve
plt.figure(figsize=(10, 6))

for model_name, model in models.items():
    # Fit the model
    model.fit(X_train, y_train)

    # Predictions
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1] # Probability estimates


    # Metrics
    accuracy = accuracy_score(y_test, y_pred) # Calculate accuracy
    cm = confusion_matrix(y_test, y_pred)
    report = classification_report(y_test, y_pred, output_dict=True)
    f1 = report['1']['f1-score']
    precision = report['1']['precision']
    sensitivity = report['1']['recall'] # also known as True Positive Rate
    specificity = cm[0][0] / (cm[0][0] + cm[0][1]) # True Negative Rate
    roc_auc = roc_auc_score(y_test, y_prob)

    # Store results with abbreviation
    short_name = model_abbreviations.get(model_name, model_name) # Get abbr
    results[short_name] = {
        'Accuracy': round(accuracy, 3),
        'Precision': round(precision, 3),
        'F1 Score': round(f1, 3),
        'Sensitivity': round(sensitivity, 3),
        'Specificity': round(specificity, 3),
        'ROC AUC': round(roc_auc, 3)
    }

    # ROC Curve
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    plt.plot(fpr, tpr, label=f'{short_name} (AUC = {roc_auc:.3f})')

# Plot ROC Curve
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend()
plt.show()

```



This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)